

At fixed T=512, batching exposes the linear-memory advantage

