

WBIC local-learning coefficient

$\beta = 1/\log n$, $n = 2048$ tokens
3 SGLD chains per model

attention

0.231

RAVEL

0.203

0.00 0.05 0.10 0.15 0.20 0.25 0.30 0.35

$$\hat{\lambda} = \beta n (\mathbb{E}_{\beta}[L] - L^*)$$

Posterior samples stay in a local basin

