

Local loss movement around the trained checkpoints

random unit directions, symmetric shells;
values are normalized to each model's best observed probe loss

validation loss above best probed point

validation loss above best probed point

-2.0 -1.5 -1.0 -0.5 0.0 0.5 1.0 1.5 2.0

distance from trained weights (% of trainable weight norm)

ravel

attention

This is a local smoothness diagnostic, not a formal RLCT estimate.

