

Attention's intermediate state grows quadratically

48 MiB

$B \times T$ is held near 1024,
so this growth is from T itself.

24 MiB

12 MiB

6 MiB

128

256

512

1024

context length T

attention score tensor MiB / forward

50
40
30
20
10
0